

# Stabilizing Cell Phone Video using Inertial Measurement Sensors

Gustav Hanning, Nicklas Forslöw, Per-Erik Forssén, Erik Ringaby, David Törnqvist, Jonas Callmer  
Department of Electrical Engineering  
Linköping University

<http://www.liu.se/forskning/foass/per-erik-forssen/VGS>

## Abstract

We present a system that rectifies and stabilizes video sequences on mobile devices with rolling-shutter cameras. The system corrects for rolling-shutter distortions using measurements from accelerometer and gyroscope sensors, and a 3D rotational distortion model. In order to obtain a stabilized video, and at the same time keep most content in view, we propose *an adaptive low-pass filter algorithm to obtain the output camera trajectory*. The accuracy of the orientation estimates has been evaluated experimentally using ground truth data from a motion capture system. We have conducted a user study, where the output from our system, implemented in iOS, has been compared to that of three other applications, as well as to the uncorrected video. The study shows that users prefer our sensor-based system.

## 1. Introduction

Most mobile video-recording devices of today make use of CMOS sensors with *rolling-shutter (RS) readout* [6]. An RS camera captures video by exposing every frame line-by-line from top to bottom. This is in contrast to a *global shutter*, where an entire frame is acquired at once.

The RS technique gives rise to image distortions in situations where either the device or the target is moving. Figure 1 shows an example of how an image is distorted when using a rolling shutter. Here, vertical lines such as the flag poles appear slanted as a result of panning the camera quickly from left to right during recording. Recording video by hand also leads to visible frame-to-frame jitter. The recorded video is perceived as “shaky” and is not very enjoyable to watch.

Since mobile video-recording devices are so common, there is an interest in correcting these types of distortions. The *inertial sensors* (accelerometers and gyroscopes) present in many of the new devices provide a new way of doing this: Using the position and/or orientation of the device, as sensed during recording, the motion induced distortions can be compensated for in a *post-processing step*.



Figure 1. An example of rolling-shutter distortion. Top: Frame from a video sequence recorded with an iPod touch. Bottom: Rectification using the 3D rotation model and inertial measurements.

### 1.1. Related Work

Early work on modeling the distortions caused by a rolling-shutter exposure is described in [7].

Rolling-shutter video has previously been rectified using image measurements. Two recent, state-of-the-art methods are described in [3, 5]. To perform rectification, we use the 3D rotational model introduced in [5], but use inertial sensor data instead of image measurements.

For stabilization we use a 3D rotation-based correction as in [13, 14], and a dynamical model derived from [17]. Differences compared to [13, 14] is the use of inertial sensor data instead of image measurements, and an adaptive

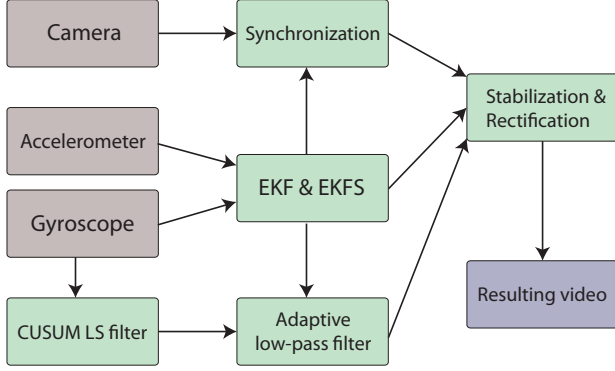


Figure 2. System overview.

### trajectory smoothing step.

How to use inertial measurement sensors for deblurring is described in [10]. This work is complementary to ours, as stabilized video may also benefit from deblurring, especially in poor light conditions. In the experiments we compare our stabilization results with results from three other stabilization applications: Deshaker [16], iMovie '11 [1], and Movie Stiller [4].

## 1.2. Overview

An overview of the stabilization and rectification system is given in figure 2. We have three data sources: camera, accelerometer and gyroscope. The data goes through a number of processing steps and the result is a stabilized video where rolling-shutter distortion has been suppressed. Each step is presented in detail in the following sections.

We have implemented this system as an iOS video stabilization application running on the 4th generation iPod touch (henceforth called the “iPod”), the iPhone 4 and the iPad 2.

## 1.3. Outline

The paper begins with a description of how the camera orientation is estimated from accelerometer and gyroscope measurements in section 2. In section 3 we present the methods for image rectification and video stabilization and deal with the problem of synchronizing a recorded video to the collected sensor data. In section 4 we evaluate the accuracy of our camera orientation estimates and compare the output of our iOS application to that of competing video stabilization software. The paper ends with some concluding remarks in section 5.

## 2. Inertial Measurement Sensors

In order to estimate the motion of the device, measurements of the current state of the system is needed. The iPod is equipped with a three-axis accelerometer measuring

acceleration  $\mathbf{a} = (a_x, a_y, a_z)^T$  and a three-axis gyroscope that measures angular velocity  $\boldsymbol{\omega} = (\omega_x, \omega_y, \omega_z)^T$ .

### 2.1. Motion Model

When modeling the motion of the device we choose to only consider the orientation. This decision is based on the results of [5] where a pure rotation model yielded the best result for rolling-shutter rectification. Our *state* is the device orientation, which we represent using a unit quaternion  $\mathbf{q}_t = (q_{0,t}, q_{1,t}, q_{2,t}, q_{3,t})^T$ . Instead of modeling  $\boldsymbol{\omega}_t$  as a state and including the measurements of the gyroscope in the measurement model,  $\boldsymbol{\omega}_t$  is used as input to the motion model. Since  $\boldsymbol{\omega}_t$  is measured and distorted by noise we introduce the measurement noise  $\mathbf{v}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q})$  and use  $\boldsymbol{\omega}_t + \mathbf{v}_t$  as angular velocity measurements. Assuming that the angular velocity is constant during sampling intervals of length  $T$  and that  $\|\boldsymbol{\omega}_{t-1}\|T$  is small, the following discrete state model can be obtained [17]:

$$\mathbf{q}_t = \left( I + \frac{T}{2} \mathbf{A}(\boldsymbol{\omega}_{t-1}) \right) \mathbf{q}_{t-1} + \frac{T}{2} \mathbf{B}(\mathbf{q}_{t-1}) \mathbf{v}_{t-1}, \quad (1)$$

where

$$\mathbf{A}(\boldsymbol{\omega}) = \begin{pmatrix} 0 & -\omega_x & -\omega_y & -\omega_z \\ \omega_x & 0 & \omega_z & -\omega_y \\ \omega_y & -\omega_z & 0 & \omega_x \\ \omega_z & \omega_y & -\omega_x & 0 \end{pmatrix}, \quad \mathbf{B}(\mathbf{q}) = \begin{pmatrix} -q_1 & -q_2 & -q_3 \\ q_0 & -q_3 & q_2 \\ q_3 & q_0 & -q_1 \\ -q_2 & q_1 & q_0 \end{pmatrix}. \quad (2)$$

### 2.2. Measurement Model

Having incorporated the gyroscope measurements into the motion model, the measurement model is only comprised of equations describing the acceleration. An accelerometer measures *proper acceleration*  $\mathbf{y}_t$ , which is the acceleration relative to free fall. The measured acceleration can be divided into two parts, gravitational acceleration  $\mathbf{g}_t^d$  and movement acceleration  $\mathbf{a}_t^d$ , both given in a frame  $d$  following the device. However, when recording a video the user usually tries to hold the camera in a steady position. Therefore we assume that only the gravitational acceleration is affecting the device by setting  $\mathbf{a}_t^d = \mathbf{0}$ . We get

$$\mathbf{y}_t = \mathbf{a}_t^d + \mathbf{g}_t^d + \mathbf{e}_t \approx \mathbf{g}_t^d + \mathbf{e}_t, \quad (3)$$

where  $\mathbf{e}_t$  is the measurement noise. In order to model the gravitational acceleration which we know is always pointing downwards to the ground we use our knowledge of the device orientation  $\mathbf{q}_t$ . The measurement equation can be written as

$$\mathbf{y}_t = \mathbf{R}(\mathbf{q}_t) \mathbf{g} + \mathbf{e}_t = \mathbf{h}(\mathbf{q}_t) + \mathbf{e}_t, \quad (4)$$

where  $\mathbf{R}(\mathbf{q}_t)$  is the rotation matrix corresponding to  $\mathbf{q}_t$  and  $\mathbf{g} = (0, 0, g)^T$  where  $g$  is the Earth’s gravity.

### 2.3. EKF Filtering

We use an *Extended Kalman Filter* (EKF) to estimate the orientation of the device and therefore, the measurement model needs to be linearized. Here, a prediction step or time update is made before incorporating the accelerometer measurements into the estimate. Denote this estimate  $\hat{\mathbf{q}}_{t|t-1}$ . Now, we can use a first order Taylor expansion around  $\hat{\mathbf{q}}_{t|t-1}$  to approximate the measurement model as

$$\mathbf{y}_t = \mathbf{h}(\mathbf{q}_t) + \mathbf{e}_t \approx \mathbf{h}(\hat{\mathbf{q}}_{t|t-1}) + \mathbf{H}(\hat{\mathbf{q}}_{t|t-1})(\mathbf{q}_t - \hat{\mathbf{q}}_{t|t-1}) + \mathbf{e}_t, \quad (5)$$

where

$$\mathbf{H}(\hat{\mathbf{q}}_{t|t-1}) = \left. \frac{\partial \mathbf{h}(\mathbf{q})}{\partial \mathbf{q}} \right|_{\mathbf{q}=\hat{\mathbf{q}}_{t|t-1}}. \quad (6)$$

To get rid of the nonlinear term  $\mathbf{h}(\hat{\mathbf{q}}_{t|t-1})$  we form a new measurement variable  $\tilde{\mathbf{y}}_t$  by rearranging (5):

$$\tilde{\mathbf{y}}_t = \mathbf{y}_t - \mathbf{h}(\hat{\mathbf{q}}_{t|t-1}) + \mathbf{H}(\hat{\mathbf{q}}_{t|t-1})\hat{\mathbf{q}}_{t|t-1} \approx \mathbf{H}(\hat{\mathbf{q}}_{t|t-1})\mathbf{q}_t + \mathbf{e}_t, \quad (7)$$

resulting in the sought linearized measurement model [17].

Assume that the noise terms,  $\mathbf{v}_{t-1}$  and  $\mathbf{e}_t$ , are zero-mean normally-distributed random variables with covariance matrices  $\mathbf{Q}_{t-1}$  and  $\mathbf{R}_t$ , respectively. Furthermore, assume  $\mathbf{q}_i \sim \mathcal{N}(\boldsymbol{\mu}_i, \mathbf{P}_i)$ , where  $\mathbf{q}_i$  is the initial state. With these assumptions the EKF can be applied to (1), (4) and (7). Note that the state  $\mathbf{q}_t$  is normalized after both measurement and time updates in order to maintain  $\mathbf{q}_t$  as a unit quaternion.

Since we use an **off-line approach** we have access to all measurements  $\mathbf{y}_t$ ,  $t = 1, \dots, N$  when estimating the orientation. We want to use this information to find the best possible state estimate  $\hat{\mathbf{q}}_{t|1:N}$ . Therefore a second filtering pass, using an *Extended Kalman Filter Smoother* (EKFS), is performed utilizing all available measurements to improve the EKF estimates. For complete derivations of the EKF and EKFS refer to [19].

### 2.4. Adaptive Smoothing using Change Detection

When capturing a movie using a handheld device it is almost impossible to hold the device steady. Due to user induced shakes and/or small movements, the resulting video will be perceived as shaky. In order to stabilize this kind of video sequence, a smooth orientation estimation is needed. One can think of this smoothed estimation as the rotational motion that the device would have had if the orientation estimate was free from small shakes or movements. A common Hanning window of length  $M$  can be used to low-pass filter the orientation. The **corresponding weights  $w_m$**  can be calculated as

$$w_m = \begin{cases} 0.5 \left( 1 + \cos \left( \frac{2\pi m}{M-1} \right) \right), & -\frac{M-1}{2} \leq m \leq \frac{M-1}{2} \\ 0, & \text{elsewhere.} \end{cases} \quad (8)$$

The window can then be applied in the time domain at time  $t$  on a window of the EKFS orientation estimates  $\hat{\mathbf{q}}_{t|1:N}$ , to obtain a stabilized orientation. The general ideas behind such a weighted averaging of quaternions can be found in [8].

When applying a window to the orientation estimates it is important to vary the window length depending on how the orientation changes. In order to suppress all user induced shakes, the window length  $M$  needs to be large. However, a large  $M$  will create an inaccurate motion for fast motion changes, **leading to large areas in the video frames that lack image information**. To solve this problem we use a *Cumulative Sum Least Squares Filter* (CUSUM LS Filter) [9] to detect when fast orientation changes occur. The idea is to divide  $\hat{\mathbf{q}}_{t|1:N}$  into different segments which can then be filtered with windows of different lengths.

The input to the CUSUM LS Filter is the norm of the angular velocity  $y_k = \|\boldsymbol{\omega}_k\|$  as measured by the gyroscope. We model this new measurement as

$$y_k = \theta_k + e_k, \quad (9)$$

where  $k = 1, \dots, N$  and  $e_k$  is the measurement noise. The LS estimation  $\hat{\theta}$  of the signal is given by

$$\hat{\theta}_k = \frac{1}{k} \sum_{i=1}^k y_i. \quad (10)$$

Define the distance measures  $s_k^1$  and  $s_k^2$  as

$$s_k^1 = y_k - \hat{\theta}_{k-1}, \quad s_k^2 = -y_k + \hat{\theta}_{k-1}, \quad (11)$$

where  $\hat{\theta}_{k-1}$  is the LS estimation based on measurements up to  $k-1$ . The distance measures will serve as input to the test statistics,  $g_k^1$  and  $g_k^2$ , defined as

$$g_k^1 = \max(g_{k-1} + s_k^1 - \nu, 0), \quad g_k^2 = \max(g_{k-1} + s_k^2 - \nu, 0), \quad (12)$$

where  $\nu$  is a small drift term subtracted at each time instance. The reason for having two test statistics is the ability to handle positive as well as negative changes in  $\|\boldsymbol{\omega}_k\|$ . Each test statistic sums up its input with the intention of giving an alarm when the sum exceeds the threshold  $h$ . **When an alarm is signaled the filter is reset and a new LS estimation starts**. The output is the index vector  $\mathbf{I}$  containing the alarm indices. These indices make up the segments  $s_i$  of  $\hat{\mathbf{q}}_{t|1:N}$ . Both  $\nu$  and  $h$  are design parameters set to achieve alarms at the desired levels.

**The adaptive smoothing algorithm uses two main filters**, both Hanning windows (8). Let these two windows be of length  $K$  and  $L$ , where  $K > L$ . First we find the big segments satisfying

$$\text{length}(s_i) > K. \quad (13)$$

In these segments  $\|\boldsymbol{\omega}_k\|$  is small or changes slowly and can thus be LP-filtered with the wider window of length  $K$ .

The remaining small segments where  $\|\omega_k\|$  changes rapidly can be filtered with the narrow window of length  $L$ . In order to achieve a smooth transition between the two windows we let the length of the windows increase or decrease with a step size  $a = 2$ . This will also guarantee that the transitioning window's symmetrical shape is kept.

### 3. Image Rectification and Video Stabilization

We use the orientation estimates obtained from the EKFS to rectify each frame of a recorded video. By low-pass filtering the quaternion estimates, we can suppress frame-to-frame jitter and stabilize the video.

#### 3.1. Camera Model

The cell phone's camera is modeled using the pinhole camera model. Using homogenous coordinates, the relations between a point in space  $\mathbf{X}$  and its projection onto the image plane  $\mathbf{x}$  are

$$\mathbf{x} = \mathbf{K}\mathbf{X} \quad \text{and} \quad \mathbf{X} = \lambda\mathbf{K}^{-1}\mathbf{x}, \quad (14)$$

where  $\mathbf{K}$  is the *camera calibration matrix* and  $\lambda$  is a non-zero scale factor.  $\mathbf{K}$  is estimated from images of a calibration pattern using the Zhang method [20], as implemented in OpenCV.

#### 3.2. Rolling Shutter Correction

If an image is captured with a rolling-shutter camera, the rows are read at different time instances and at each of these instances the camera has a certain orientation. We want to transform the image to make it look as if all rows were captured at once.

To each point  $\mathbf{x} = (x, y, z)^T$  in a frame we associate a time instance  $t(\mathbf{x})$  which can be calculated as

$$t(\mathbf{x}) = t_f + t_r \frac{y}{zN}, \quad (15)$$

where  $t_f$  is the instance in time when the image acquisition started and  $t_r$  is the readout time of the camera, which is assumed to be known.  $N$  is the number of rows in the image. The readout time,  $t_r$ , can be estimated by recording a flashing light source with a known frequency [5, 7].

The orientation of the camera at time  $t(\mathbf{x})$  is found by interpolating from the EKFS estimates. Let  $\mathbf{q}_1$  and  $\mathbf{q}_2$  be two consecutive estimates belonging to times  $t_1$  and  $t_2$  such that  $t_1 \leq t(\mathbf{x}) \leq t_2$ . We use spherical interpolation [15] to find the camera orientation:

$$\mathbf{q}(\mathbf{x}) = \text{SLERP}(\mathbf{q}_1, \mathbf{q}_2, \frac{t(\mathbf{x}) - t_1}{t_2 - t_1}). \quad (16)$$

Now let  $\mathbf{q}_m$  be the camera orientation for the middle row of the image. The rotation from the middle row to the row containing  $\mathbf{x}$  is

$$\mathbf{q}_m(\mathbf{x}) = \mathbf{q}_m^{-1} \odot \mathbf{q}(\mathbf{x}), \quad (17)$$

where  $\odot$  denotes quaternion multiplication. For each point in the image we can calculate its position during the acquisition of the middle row of the frame.

$$\mathbf{X} = \lambda\mathbf{K}^{-1}\mathbf{x} \quad (18)$$

$$\begin{pmatrix} 0 \\ \mathbf{X}' \end{pmatrix} = \mathbf{q}_m(\mathbf{x})^{-1} \odot \begin{pmatrix} 0 \\ \mathbf{X} \end{pmatrix} \odot \mathbf{q}_m(\mathbf{x}) \quad (19)$$

$$\mathbf{x}' = \mathbf{K}\mathbf{X}' \quad (20)$$

$\mathbf{x}'$  is thus the new, rectified position. Using this technique, all rows are aligned to the middle row. One can get another reference row by replacing  $\mathbf{q}_m$  in (17).

#### 3.3. Video Stabilization

If a video is recorded by hand, it will be shaky and the EKFS estimates will vary a lot from frame to frame. To remove these quick changes in orientation, i.e. to stabilize the video, **we low-pass filter the quaternions as described in section 2.4.**

For each frame we find the low-pass filtered camera orientation corresponding to the middle row,  $\mathbf{q}_m^{LP}$ . The stabilized location of an image point is then calculated as

$$\mathbf{q}_\Delta = \mathbf{q}_m^{-1} \odot \mathbf{q}_m^{LP} \quad (21)$$

$$\begin{pmatrix} 0 \\ \mathbf{X}'' \end{pmatrix} = \mathbf{q}_\Delta \odot \begin{pmatrix} 0 \\ \mathbf{X}' \end{pmatrix} \odot \mathbf{q}_\Delta^{-1} \quad (22)$$

$$\mathbf{x}'' = \mathbf{K}\mathbf{X}''. \quad (23)$$

#### 3.4. Synchronization

Synchronizing the sensor data to the recorded video is important as failure to do so leads to a poor result, often worse than the original video. **Both sensor data and frames are time stamped on the iPod, but there is an unknown delay between the two.** We add a term  $t_d$  to (15), representing this delay:

$$t(\mathbf{x}) = t_f + t_d + t_r \frac{y}{zN}. \quad (24)$$

To find  $t_d$ , a number of points are randomly selected in each of  $M$  consecutive frames of the video (Harris points give similar results). The points of a frame are tracked to the next frame by the use of the KLT tracker [12] implemented in OpenCV. They are then re-tracked to the original frame and only those that return to their original position, within a threshold of 0.5 pixels, are kept. Figure 3 shows 25 random points in a video frame. Successfully tracked points are marked with yellow circles. Points for which the tracking failed are marked with red circles.

For each resulting point correspondence  $\mathbf{x} \leftrightarrow \mathbf{y}$ , we rotate  $\mathbf{y} = (u, v, w)^T$  back to the position it had at time instance  $t(\mathbf{x})$ :

$$\mathbf{Y} = \lambda\mathbf{K}^{-1}\mathbf{y} \quad (25)$$








Figure 3. Twenty-five random points in a frame. The points are tracked to the next frame in the video, and then re-tracked. Yellow circles mark those points that returned to their original position. Red circles mark points for which the tracking failed.

$$\begin{pmatrix} 0 \\ \mathbf{Y}' \end{pmatrix} = \mathbf{q}(\mathbf{x}) \odot \mathbf{q}(\mathbf{y})^{-1} \odot \begin{pmatrix} 0 \\ \mathbf{Y} \end{pmatrix} \odot \mathbf{q}(\mathbf{y}) \odot \mathbf{q}(\mathbf{x})^{-1} \quad (26)$$

$$\mathbf{y}' = \mathbf{K}\mathbf{Y}'. \quad (27)$$

$\mathbf{q}(\mathbf{x})$  and  $\mathbf{q}(\mathbf{y})$  are the orientations of the camera at time instances  $t(\mathbf{x})$  and  $t(\mathbf{y})$ , acquired by interpolation from the EKFS estimates as in section 3.2. A cost function  $J$  can now be defined as

$$J = \sum_k d(\mathbf{x}_k, \mathbf{y}'_k)^2, \quad (28)$$

where  $d(\mathbf{x}, \mathbf{y})$  is the pixel distance between  $\mathbf{x}$  and  $\mathbf{y}$ ,

$$d(\mathbf{x}, \mathbf{y})^2 = \left(\frac{x}{z} - \frac{u}{w}\right)^2 + \left(\frac{y}{z} - \frac{v}{w}\right)^2. \quad (29)$$

Since  $J$  could possibly have local minima, we use a grid search to find the global minimum. Figure 4 shows  $J$  as a function of  $t_d$ , for a sample video. As seen, there is a distinct minimum at  $t_d = 0.144$  and two other, local, minima.




Figure 4. The cost function  $J$ , evaluated at different values of  $t_d$ .

## 4. Experiments

The orientation estimates from the EKF and EKFS have been evaluated on experimental data using ground truth estimates from a motion capture system. In addition, a user study was conducted comparing our iOS application output videos to those of competing software.




Figure 5. Rig used during the motion capture experiments. The reflective balls are used to track the orientation with precision on a sub degree level.

### 4.1. Orientation Accuracy

To evaluate the quality of the orientation estimates a series of experiments were conducted using a Vicon MX motion capture system [18]. Figure 5 displays the rig used during the experiments.

The reflective balls attached to the wooden plate are tracked by 10 infrared cameras and infrared lamps placed on the walls of the 3 by 5 meter room. The sampling frequency of the system is 250 Hz and the precision of the estimated orientation is on a sub degree level. In our experiments the iPod gyroscope and accelerometer sampling frequencies were clocked at approximately 60 and 95 Hz respectively. Note that the sampling frequencies for both sensors are non-constant.

In the evaluation two data sets were used. Both sets were bias corrected using estimates from data where the sensors were at rest. In order to estimate the accelerometer bias we used 14 different data sets where the accelerometer was positioned in different orientations. For each set we calculate the mean acceleration and then fit a sphere to these points. The bias can be obtained as the translation of the sphere relative to the origin. In addition we scale the sphere to obtain the acceleration scaling factor since we know that its radius should be equal to  $g$ . The first Vicon data set was used to tune the covariance matrices  $\mathbf{Q}$  and  $\mathbf{R}$  used in the EKF. Since the Vicon data is considered ground truth we used this data to set the initial orientation  $\mathbf{q}_i$ . The initial state covariance was set as  $\mathbf{P}_i = 10^{-10}\mathbf{I}$ .

For the second experiment, data was logged having the wooden plate parallel to the ground. The person holding the plate rotates to the left and back. In order to better analyze the results, the quaternions are converted to Euler angles using the aerospace sequence, i.e. ZYX [11]. Using this representation the orientation can be considered a result of three composite rotations by the angles  $\psi$ ,  $\theta$  and  $\phi$ . Figure 6 displays the EKF and EKFS estimated Euler angles together with ground truth data from the Vicon system. As can be

seen, the estimates are accurate with only a few degrees offset in each angle. Note that it is impossible maintain accurate global estimates over a long period of time. This is because we lack an additional heading reference.




Figure 6. Estimated Euler angles  $\psi$ ,  $\theta$  and  $\phi$  using the EKF (solid) and EKFS (dashed) plotted together with ground truth data from the Vicon system (dash dot).

The top plot in Figure 7 displays the global angular error of the EKF and EKFS estimates. The image rectification algorithm relies on good estimates of the relative orientation change from the capture time instance of the first image row to the last. The bottom plot in Figure 7 displays the angular error of this relative orientation change using a time interval of 30 ms corresponding to the readout time of the iPod.




Figure 7. The top plot shows the global angular error. At the bottom the relative angular error is displayed. EKF (solid), EKFS (dashed).

Here, at samples corresponding to beginning and end of

the rotational motion, the error is relatively large (approximately samples 90, 190, 275, 360). This is mainly caused by gyroscope scaling errors leading to larger angular errors when the orientation changes fast.

Table 1 shows the mean angular errors of the global and relative estimates together with errors from pure integration of the gyroscope measurements. As can be seen, the EKF reduces the global error compared to the pure integration. The EKFS overall improves the EKF estimates although only moderately. For the relative estimates we only see a small improvement using the EKF and EKFS.

method	glob. ang. mean error	rel. ang. mean error
Gyro	4.3652°	0.1123°
EKF	3.5455°	0.1104°
EKFS	3.5378°	0.1091°

Table 1. The mean angular error of the global and relative estimates shown in Figure 7 together with the angular errors from pure integration of the gyroscope measurements.

To summarize, the estimates are accurate although the improvements on the EKF estimates using the EKFS are small. However, normally there is an uncertainty in the initial state  $\mathbf{q}_i$  which leads to poor estimates for the first samples using only the EKF. The backwards EKFS pass greatly improves these estimates.

## 4.2. User Study

A user study was conducted to evaluate the proposed video stabilization and rectification system. Two video sequences, approximately ten seconds long, were recorded with the iPod. The sequences were then processed by four different applications: Deshaker [16], iMovie '11 [1], Movie Stiller [4] and DollyCam (our iOS application).

The first video was captured while walking. This clip has lots of shaking and a number of moving objects (people). The second video is similar to the first but also contains fast panning to get some noticeable rolling-shutter distortion. It features fewer moving objects. The videos can be found in the *supplemental material*, together with our results.

The settings for each application were chosen in a way that let them use their particular strengths. For example, if an application supports some kind of edge-filling, one can usually allow a little less zoom and still get good results. The zoom values were chosen high enough to remove most of the black edges, but not excessively high since this causes the videos to lose sharpness.

### 4.2.1 Deshaker

Deshaker is a video stabilizer plugin for the open source video processing utility VirtualDub [2]. In addition to stabilization, Deshaker can also do rolling-shutter correction

and supports edge compensation where information from previous and future frames are used.

Deshaker version 2.7 was used to process the videos and most of the settings were left at their default values. The most important parameter changes were setting the rolling-shutter “amount” to 72% (the video was recorded at 24 fps) and enabling edge compensation. A fixed zoom of 10% was used for the first video sequence and 20% fixed zoom was used for the second. This was to avoid too large borders.

#### 4.2.2 iMovie ’11

iMovie ’11 is a video editing software created by Apple. iMovie performs video stabilization and rolling-shutter correction but does not support any edge-filling. The amount of stabilization is chosen by setting a maximum zoom value. The amount of rolling-shutter correction can be chosen as “None”, “Low”, “Medium”, “High” or “Extra High”. A zoom of 27% and Extra High rolling-shutter correction were used for both video clips.

#### 4.2.3 Movie Stiller

Movie Stiller is an application for the iPhone, iPod and iPad that can stabilize movies recorded on the device. In contrast to the other application used in the user study, it does not correct for rolling-shutter distortions. There are two parameters that a user of Movie Stiller can choose: Stabilization Strength and Default Scale. Setting a large stabilization strength without increasing the scale would then lead to black borders in the resulting video, because Movie Stiller does not support any edge-filling.

A zoom value of 27% was used in both videos and the stabilization strength was tuned so that little or no borders were visible in the output videos. The stabilization strength was set to 3.5 and 0.3, respectively. Version 1.1 of Movie Stiller was used to do the stabilization.

#### 4.2.4 Our iOS Application

For stabilization, our iOS application (named DollyCam) used an adaptive low-pass filter with minimum filter size  $L = 29$  and maximum filter size  $K = 99$  (as described in section 2.4). The zoom was set to 15% and 25%, respectively. Extrapolation was used to fill any black borders.

#### 4.2.5 Method

The user study was conducted as a blind experiment, where users were shown pairs of videos and asked to choose the one they thought looked the best. They had no prior knowledge about what software was used or even what the applications tried to achieve (*i.e.* a stable video without rolling-shutter distortion).




Figure 8. Total number of points for each application (DollyCam is our application). Dark blue: first video. Light blue: second video.

There were five versions of each of the two recorded video sequences: the original recording and the output of the four applications described above. All versions were compared to each other, meaning that a user would make a total of 20 decisions for the two movies. The order in which the versions appeared was randomized, but was the same for all users.

#### 4.2.6 Results

30 people, aged 18 to 36, participated in the user study. The applications were awarded one point for each “win” against another application. Figure 8 shows the total number of points that each application received for the first and second video sequence, marked with dark and light blue, respectively.

As seen, our iOS application got the highest number of points for the first video, followed by Deshaker. iMovie ’11 and Movie Stiller performed similarly, and the original video had the least number of wins. Our iOS application got 106 out of 120 points possible, losing 9 points to Deshaker (but won 21), 3 points to the original video, and 1 point each to iMovie ’11 and Movie Stiller.

From the data we draw the conclusion that it is certainly possible to improve upon the original video sequence and that in general, people preferred the stabilized versions to the original one. Two of the users did not agree though, ranking the original video first. This is possibly because it was the only one not suffering from the loss in image quality caused by zooming.

Our iOS application scored the highest number of points also for the second sequence, with Deshaker not far behind. Despite the lack of rolling-shutter correction, Movie Stiller got the third most points. The original video sequence outperformed iMovie ’11, which got the least amount of points. Our iOS application collected exactly 100 points, losing 14 points to Deshaker, and 6 points to the original video.

The increased number of points for the original video (in comparison to the first video sequence) indicates that this

second video sequence was more difficult to stabilize and rectify. Still, our iOS application and Deshaker performed well, whereas iMovie '11 failed to improve upon the original video. 5 people thought that none of the applications gave a better result than the original video, ranking it the highest.

To sum up, the methods of video stabilization and rolling-shutter correction presented in this paper performs equal to or better than the algorithms implemented in Deshaker, iMovie '11 and Movie Stiller, at least when applied to the two video sequences used in the study. We can also see that the output of our iOS application clearly improves over the original video sequences. One should remember though, that our iOS application has access to data that the other applications do not, namely the accelerometer and gyroscope measurements.

## 5. Concluding Remarks

In this paper, we have presented video stabilization and image rectification methods that utilize sensor data from accelerometer and gyroscope. The methods have been implemented as an iOS application, DollyCam, that is available in the Apple app store. The application has been compared to other video stabilization software in a user study, showing that this sensor-based approach to stabilization and rectification produces results better than those of Deshaker, iMovie '11 and Movie Stiller.

A disadvantage with the proposed system is that only recent mobile devices are equipped with the needed inertial sensors. The sensors need to have high enough update rates and accuracy, so that orientation estimates of reasonable quality can be extracted. Also, using sensors introduces the problem of synchronizing the recorded video sequences with the sensor data. Using the wrong synchronization offset to the gyroscope and accelerometer timestamps gives a very poor result, so this issue must be treated with care.

The use of inertial sensors for device motion estimation greatly decreases computation time, compared to image processing. This is of course important when working with mobile devices with limited processing power. As an example, the filtering computation time of a one minute video is roughly one second. Total running time on the iPhone 4 is about  $3\times$  the length of a clip at 720p resolution.

Another advantage with inertial measurements compared to image processing is that the device motion can be accurately recovered, regardless of what is imaged. This means that scenes with, for example, many moving objects or poor light conditions are no longer problematic.

## Acknowledgements

This research has been funded by the CENIT organisation at the Linköping Institute of technology, and by

Linköping University. Use of the Vicon real-time tracking system was facilitated by the UAS Technologies Lab, Artificial Intelligence and Integrated Computer Systems Division (AIICS) at the Department of Computer and Information Science (IDA). <http://www.ida.liu.se/divisions/aiics/aiicssite/index.en.shtml>

## References

- [1] Apple Inc. iMovie'11 video stabilizer, 2011. <http://www.apple.com/ilife/imovie/>. 2, 6
- [2] Avery Lee. VirtualDub video capture/processing utility, 2011. <http://www.virtualdub.org/>. 6
- [3] S. Baker, E. Bennett, S. B. Kang, and R. Szeliski. Removing rolling shutter wobble. In *IEEE CVPR'10*, June 2010. 1
- [4] Creaceed. Movie Stiller video stabilizer, 2011. <http://www.creaceed.com/elasty/iphone/>. 2, 6
- [5] P.-E. Forssén and E. Ringaby. Rectifying rolling shutter video from hand-held devices. In *IEEE CVPR'10*, June 2010. 1, 2, 4
- [6] A. E. Gamal and H. Eltoukhy. CMOS image sensors. *IEEE Circuits and Devices Magazine*, May/June 2005. 1
- [7] C. Geyer, M. Meingast, and S. Sastry. Geometric models of rolling-shutter cameras. In *6th OmniVis WS*, 2005. 1, 4
- [8] C. Gramkow. On averaging rotations. *International Journal of Computer Vision*, 42(1/2):7–16, 2001. 3
- [9] F. Gustafsson. *Adaptive Filtering and Change Detection*. John Wiley & Sons, ISBN: 978-0-471-49287-0, 2000. 3
- [10] N. Joshi, S. B. Kang, C. L. Zitnick, and R. Szeliski. Image deblurring using inertial measurement sensors. In *Proceedings of ACM Siggraph*, 2010. 2
- [11] J. B. Kuipers. *Quaternions And Rotation Sequences, A Primer With Applications To Orbits, Aerospace, And Virtual Reality*. Princeton University Press, 2nd edition, 1999. 5
- [12] B. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision. In *IJCAI81*, pages 674–679, 1981. 4
- [13] C. Morimoto and R. Chellappa. Fast 3D stabilization and mosaic construction. In *CVPR97*, pages 660–665, San Juan, Puerto Rico, June 1997. 1
- [14] E. Ringaby and P.-E. Forssén. Efficient video rectification and stabilisation for cell-phones. *IJCV*, on-line June 2011. 1
- [15] K. Shoemake. Animating rotation with quaternion curves. In *Int. Conf. on CGIT*, pages 245–254, 1985. 4
- [16] G. Thalin. Deshaker video stabilizer plugin v2.7 for VirtualDub, 2011. <http://www.guthspot.se/video/deshaker.htm>. 2, 6
- [17] D. Törnqvist. *Estimation and Detection with Applications to Navigation*. PhD thesis, Linköping University, 2008. 1, 2, 3
- [18] VICON - Motion capture systems from VICON, 2011. <http://www.vicon.com/>. 5
- [19] B. M. Yu, K. V. Shenoy, and M. Sahani. Derivation of Extended Kalman Filtering and Smoothing Equations, 2004. [http://www.ece.cmu.edu/~byronyu/papers/derive\\_eks.pdf](http://www.ece.cmu.edu/~byronyu/papers/derive_eks.pdf). 3
- [20] Z. Zhang. A flexible new technique for camera calibration. *IEEE TPAMI*, 22(11):1330–1334, 2000. 4